



The 8th KSCSE AI Winter School @ High1 Resort



Using Multiple GPUs for Distributed Deep Learning on Neuron

Soonwook Hwang

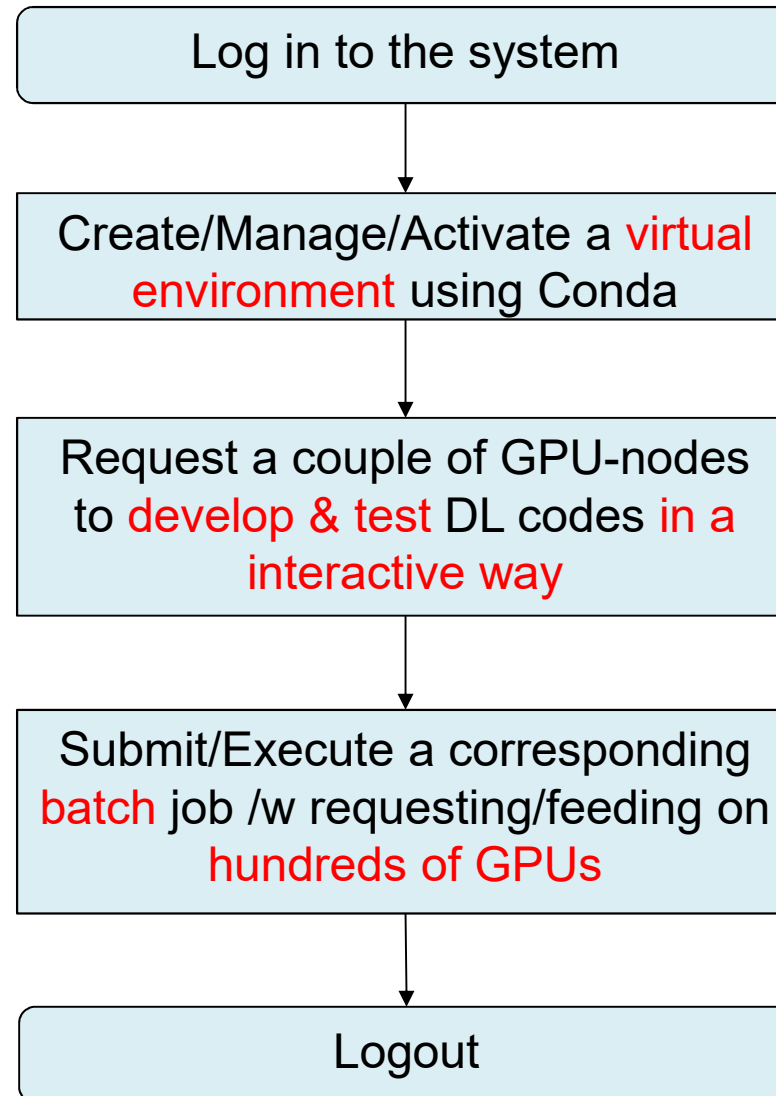
December 21, 2022

Korea Institute of Science and Technology Information

Some Motivational thoughts on large-scale DL/LM on top of national supercomputing facilities

- **KISTI-6 /w ~600PF is coming in 2 years**
 - ✓ It is expected that several thousands of GPUs(?) are available
- **Is large-scale LM (Language Model) training an exclusive task that can be conducted by big tech companies (e.g., Google, Meta, MS) running data center facilities?**
- **Why is it that large-scale LM R&D is so hard in Korea?**
 - ✓ Lack of computing resources
 - ✓ Lack of datasets
 - ✓ Lack of skills??
- **What can KISTI do in contributing to large-scale LM R&D in Korea?**
 - ✓ KISTI is uniquely positioned to running a general-purpose national supercomputing facility in Korea
- **Is KISTI's supercomputing facility easy to access for users to do large-scale distributed deep learning R&D?**
- **What are the most significant barriers that prevent users from having access to KISTI supercomputing facilities in conducting large-scale distributed training?**
 - ✓ Is it because of the traditional batch-scheduling based service?
 - ✓ ??

Distributed Training Common Practices/Routines on Supercomputers



Objectives



- **Guide users to run his/her DL codes using multiple GPU nodes on Neuron (KISTI GPU Cluster)**
- **Introduce how to set up a Horovod-enabled virtual environment**
 - install Conda and create/set up his/her virtual environment
 - install Horovod on his/her virtual environment
- **Introduce how to run Distributed DL training on Neuron**
 - Horovod
 - Singularity Container
 - ...

Agenda

- **Why distributed training?**
- **Distributed DL Approaches**
 - Data Parallel
 - Parameter Server(s)
 - Ring AllReduce
 - Model Parallel
 - Pipeline parallel
 - Tensor parallel
 - Multi-dimensional (3D) Parallel
- **KISTI GPU Cluster: Neuron**
- **Hands-on Exercises**
 - Conda Virtual Environment
 - Distributed Data Parallel(DDP) using Horovod
 - DDP using Singularity Container

Deep Learning in a slide

Courtesy of Prof. Sanjeev Arora



θ : parameters

(X, Y) : (Input, Label) pair of
Training data

$\ell(\theta, x, y)$

Loss function (how well net output **matched true label** y on point x) Can be l_2 , cross-entropy....

Objective $\operatorname{argmin}_{\theta} E_i[\ell(\theta, x_i, y_i)]$


Gradient Descent $\theta^{(t+1)} \leftarrow \theta^{(t)} - \eta \nabla_{\theta} (E_i[\ell(\theta^{(t)}, x_i, y_i)])$

Stochastic GD: Estimate ∇ via small sample of training data.

Why Distributed Training?

Why Distributed Training


- Model size is growing too big to fit
 - GPT-1('18): 110M, GPT-2('19):1.5B GPT-3('20):175B



<https://medium.com/riselab/ai- and- memory- wall- 2cb4265cb0b8>

AI and Memory Wall


- The amount of memory required to train different DL models



<https://medium.com/riselab/ai-and-memory-wall-2cb4265cb0b8>

Scaling Laws

- **Larger models are significantly more sample-efficient**
 - ✓ optimally compute-efficient training involves training very large models on a relatively modest amount of data and stopping significantly before convergence.



Scaling Laws for Neural Language Models

<https://arxiv.org/pdf/2001.08361.pdf>

Why Distributed Training

- **With More compute power and memory using multiple devices (GPU/TPU/CPU)**
 - Enable training large model
 - Speed up model training and shorten training time
 - Shorter training time allows you to do more experiments to reach your modeling goal

Scaling with Distributed Training

- (Tensorflow) Training ResNet50 with ImageNet data on Neuron
 - /apps/applications/singularity_images/imagenet (~150G)
 - train #: ~1,280,000, validation #: 50,000, test #: 100,000 images
 - GPU #8: **10.1 min/epoch**, GPU #4: **17.6 min**, GPU #1: **134.7min**

```
[gpu37] $ mpirun -np 8 python tf_keras_imagenet_resnet50.py
Epoch 1/90
5004/5004 [=====] - 604s 118ms/step - loss: 6.5324 -
accuracy: 0.0560 - top_k_categorical_accuracy: 0.1529
```

```
[gpu37] $ mpirun -np 4 python tf_keras_imagenet_resnet50.py
Epoch 1/90
10009/10009 [=====] - 1054s 104ms/step - loss:
6.3358 - accuracy: 0.0678 - top_k_categorical_accuracy: 0.1794
```

```
[gpu37] $ mpirun -np 1 python tf_keras_imagenet_resnet50.py
Epoch 1/90
40037/40037 [=====] - 8081s 202ms/step - loss:
6.0712 - accuracy: 0.0817 - top_k_categorical_accuracy: 0.2053
```

Scaling with Distributed Training

- (Pytorch) Training ResNet50 with ImageNet data on Neuron
 - /apps/applications/singularity_images/imagenet (~150G)
 - train #: ~1,280,000, validation #: 50,000, test #: 100,000 images
 - GPU #8: 15.04 min/epoch, GPU #4: 28.19, GPU #1: 3:11:30

```
[gpu37] $ mpirun -np 8 python pytorch_imagenet_resnet50.py
Train Epoch #1: 100%|██████████| 5005/5005 [15:26<00:00, 5.40it/s, loss=5.63, accuracy=5.33]
```

```
[gpu37] $ mpirun -np 4 python pytorch_imagenet_resnet50.py
Train Epoch #1: 100%|██████████| 10010/10010 [28:19<00:00, 5.89it/s, loss=5.53, accuracy=6.03]
```


```
[gpu37] $ mpirun -np 1 python pytorch_imagenet_resnet50.py
Train Epoch #1: 100%|██████████| 40037/40037 [3:11:30<00:00, 3.48it/s, loss=5.22, accuracy=8.81]
```

```
[gpu37] $ mpirun -np 1 python pytorch_imagenet_resnet50.py -fp16-allreduce
Train Epoch #1: 100%|██████████| 40037/40037 [2:01:07<00:00, 5.51it/s, loss=5.22, accuracy=8.81]
```


Distributed DL Approaches

Distributed DL approaches

Model Parallelism

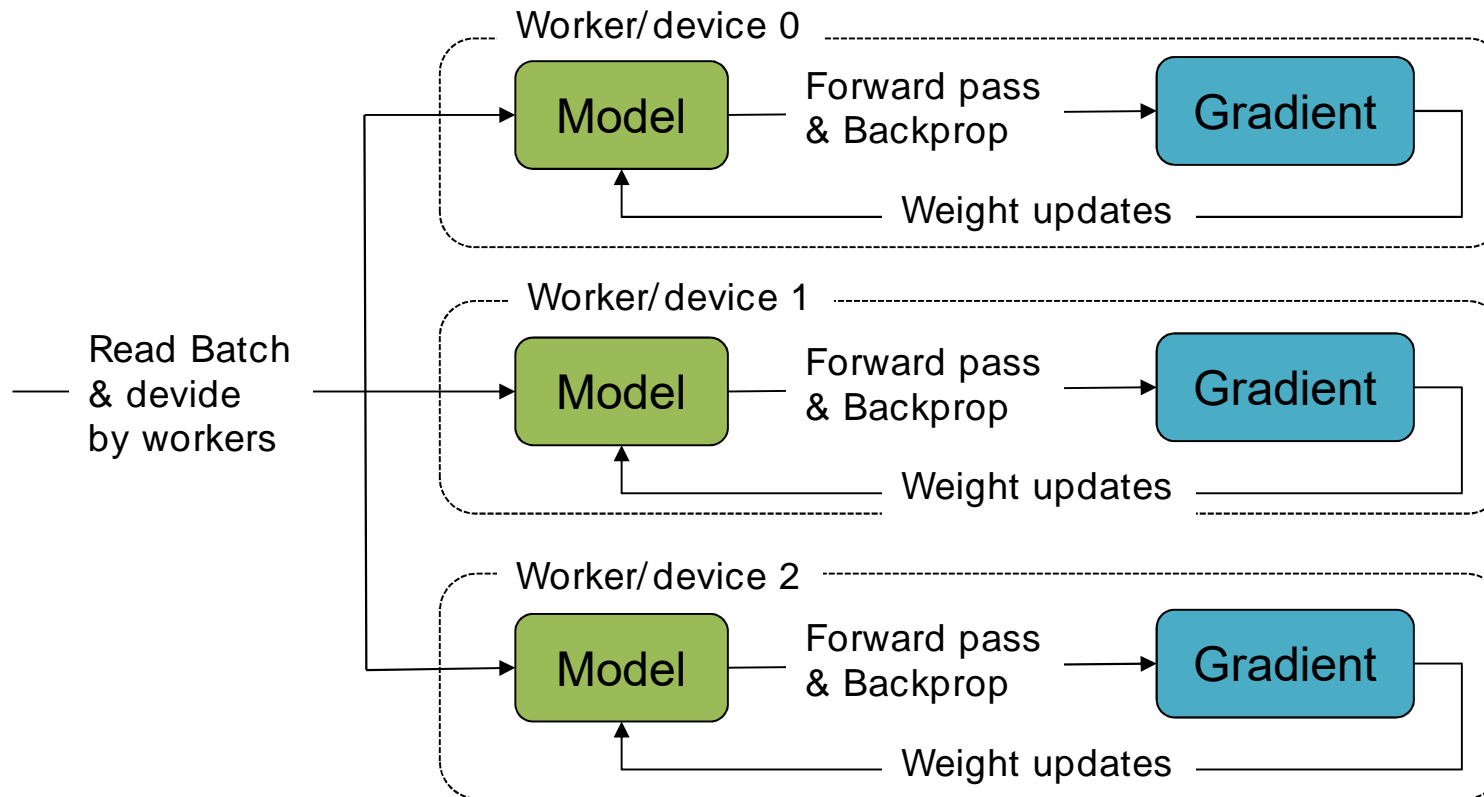


Data Parallelism



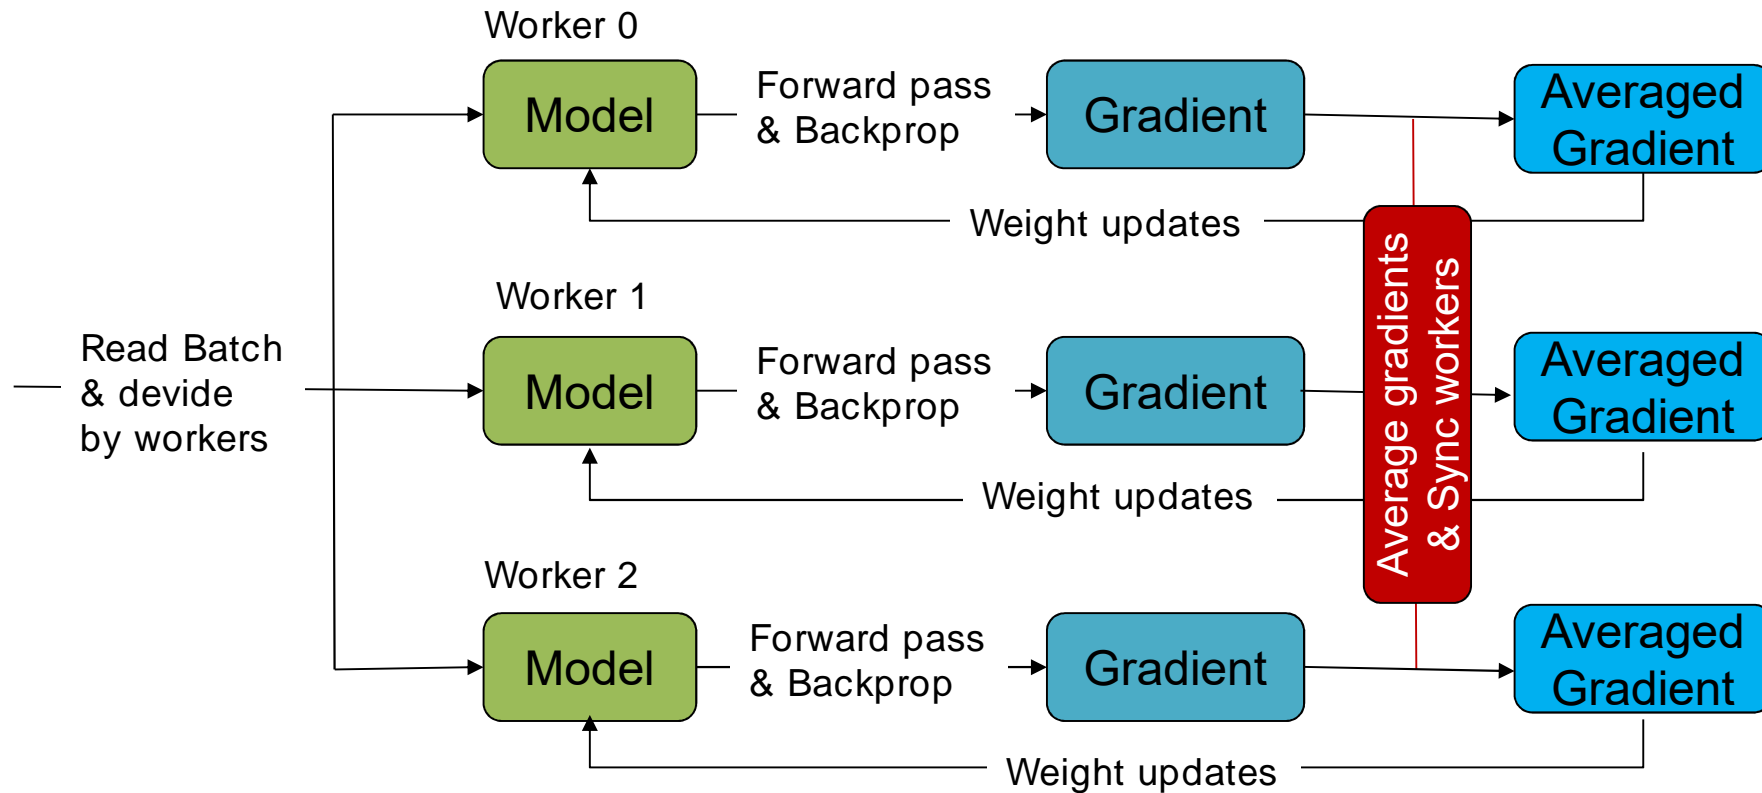
- Different parts of **model** running on multiple GPUs
- Model is too large, which cannot fit in a single device
- Different parts of **data** running on multiple GPUs
- Data is too large, which need to be processed in parallel

Data Parallel w/o Sync




- How to divide/distribute batches to workers
- When/How to synchronize workers
- How to maintain consistency in states

Data Parallel



Data Parallel




- 1. CPU feeding mini- batches to each worker**
- 2. Each worker sending its own gradient to rank 0 worker**
- 3. Rank 0 worker broadcasting the average gradient to each workers**

How to synchronize model parameters

- **Parameter Server (centralized)**
 - Synchronous
 - Asynchronous
- **Sync Allreduce (decentralized)**

Parameter Server (PS)


1. **Split the training data into shards** and assign a model replica to each data shard
2. For each model replica, fetch the parameters from the centralized **sharded parameter server**
3. Gradients are computed per model and pushed back to the parameter server



Each data shard stores a subset of the complete training data.

Asynchronous vs. Synchronous PS

Downpour SGD:
*Online Asynchronous Stochastic
Gradient Descent*



Sandblaset L- BFGS:
*Batch Distributed Parameter Storage
and Manipulation*


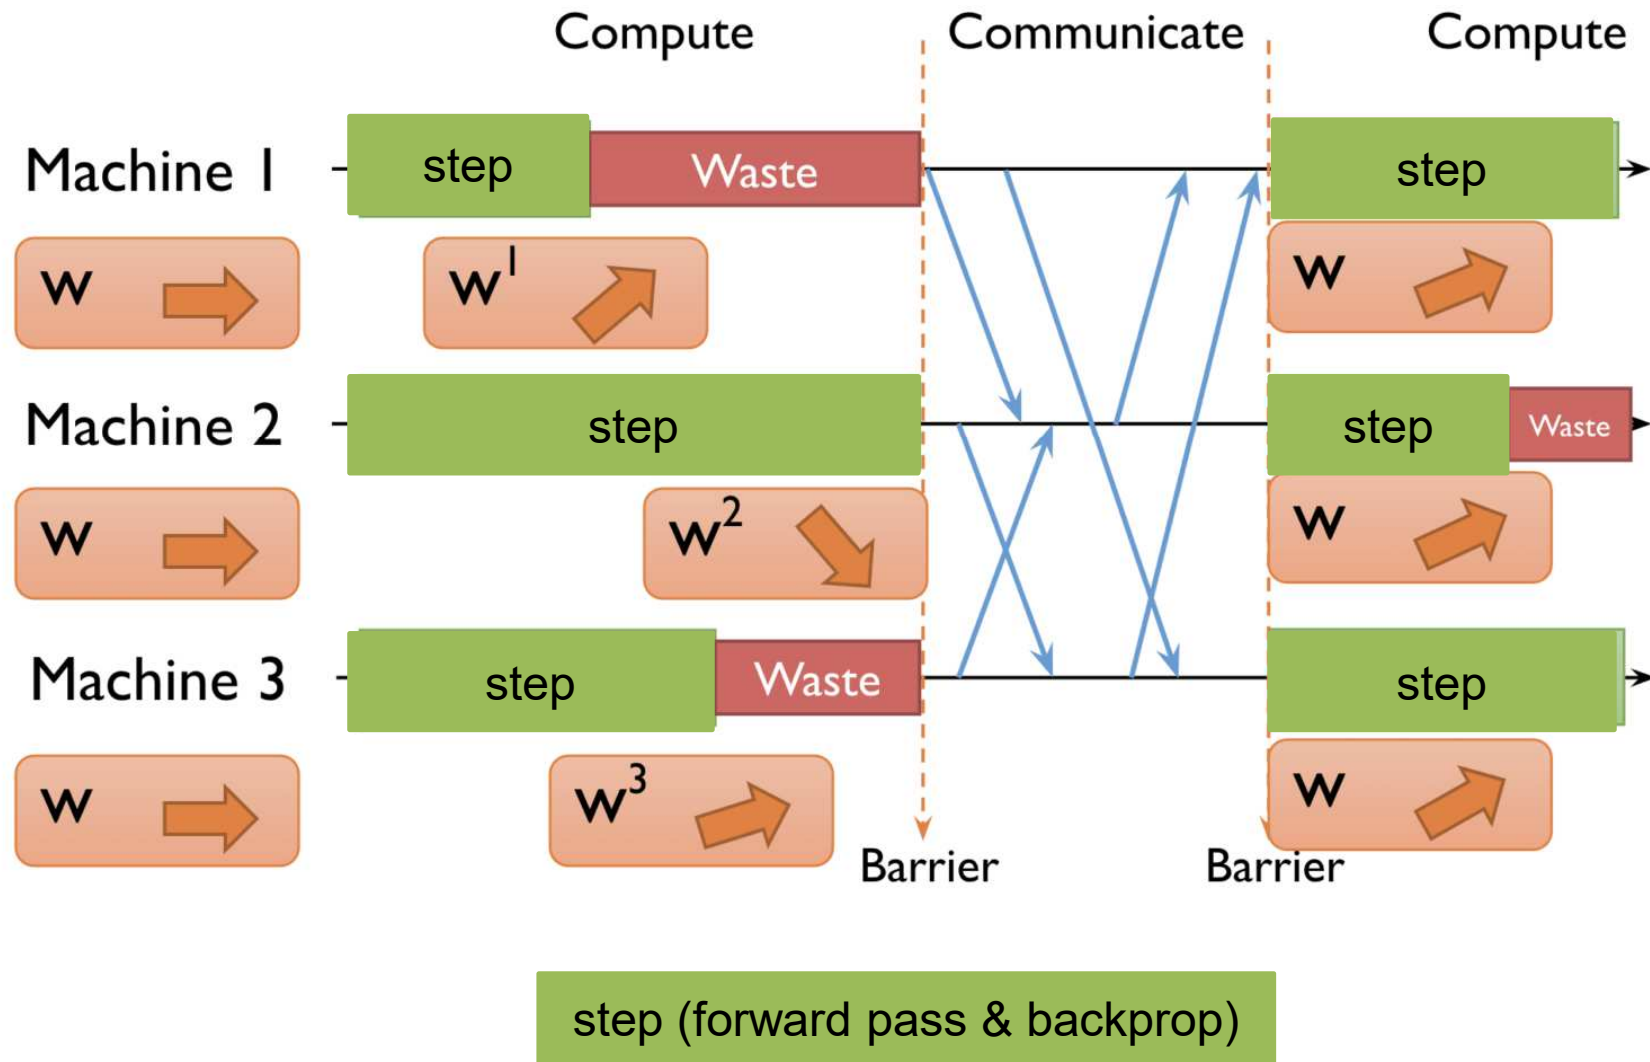


Figure 2: Left: Downpour SGD. Model replicas asynchronously fetch parameters w and push gradients Δw to the parameter server. Right: Sandblaster L-BFGS. A single 'coordinator' sends small messages to replicas and the parameter server to orchestrate batch optimization.

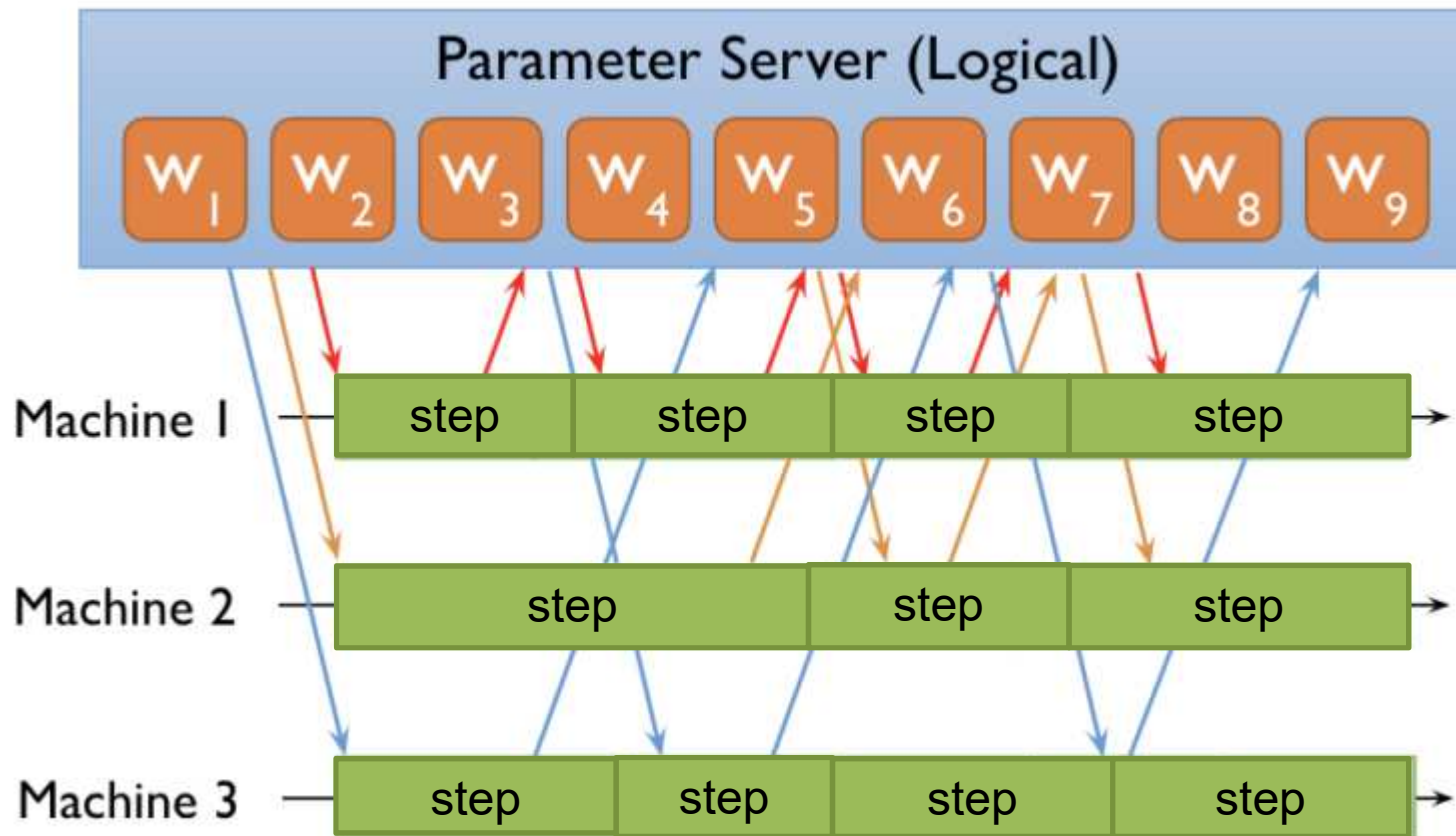
Figure from Large Scale Distributed Deep Networks

<https://static.googleusercontent.com/media/research.google.com/ko//archive/large-deep-networks-nips2012.pdf>

Synchronous PS



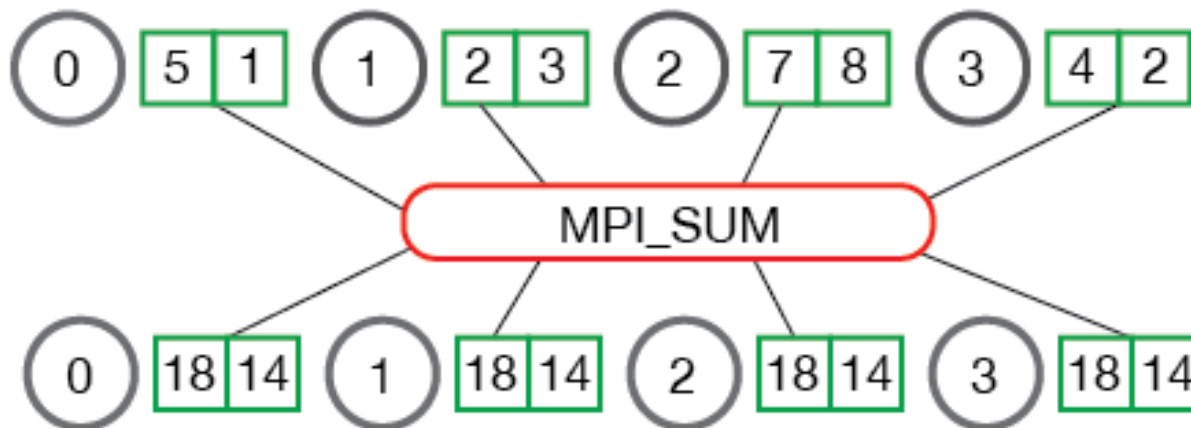
Asynchronous PS



Allreduce operation

- a collective communication operation that reduce a set of arrays on distributed workers to a single array that is then re-distributed back to each workers

MPI_Allreduce



Ring Allreduce

- Synchronized w/o parameter server(s)

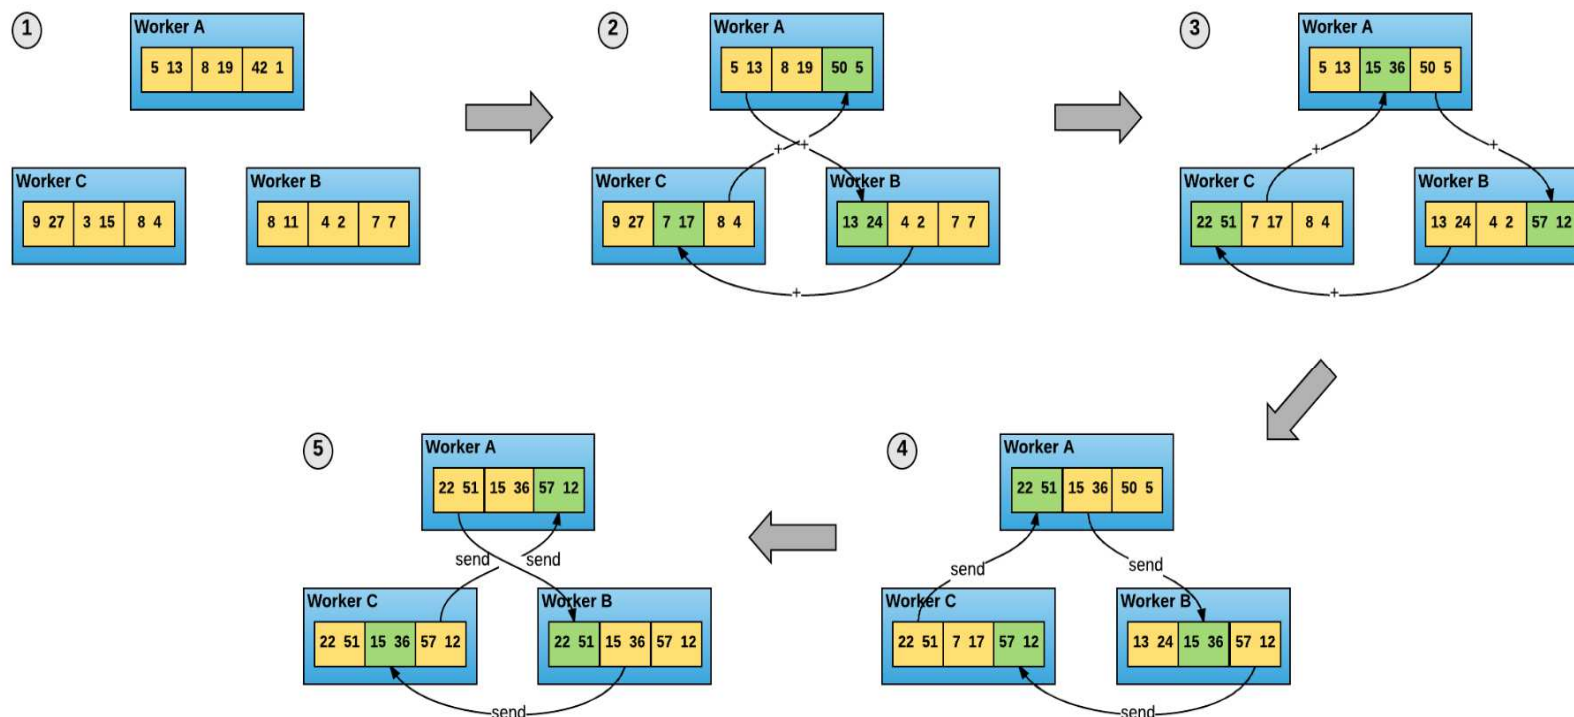


Figure 4: The ring-allreduce algorithm allows worker nodes to average gradients and disperse them to all nodes without the need for a parameter server.

Horovod: fast and easy distributed deep learning in Tensorflow

<https://arxiv.org/pdf/1802.05799.pdf>

Each step for each worker in data parallel

- 1) Each GPU performs the **forward pass** on a **different slice of the input data** (mini-batch) to **compute the loss**
- 2) Each GPU **computes the gradients** based on the **loss function**
- 3) These gradients are **aggregated across all of the devices, via an Allreduce operation**
- 4) The optimizer **updates the weights using the average gradients**, thereby keeping the devices in sync

horovod code snippet using tensorflow


```
training_step(images, labels):  
    with tf.GradientTape() as tape:  
        # 1&2) compute the gradients  
        probs = mnist_model(images,)   
        loss_val = loss(labels, probs)
```

```
# adding Horovod distributed gradients  
tape = hvd.DistributedGradientTape(..)
```

```
#3) aggregating all the gradients  
grads = tape.gradient(loss_val,...)
```

```
#4) updating the weights  
opt.apply_gradients(grads, )
```

Latency Breakdown of Data Parallel




- **COMM is very expensive**
- **What if averaging model parameters after OPT (optimizer step) rather than averaging gradients?**
- **Is there a way to reduce latency and make it fast?**

Model Parallel

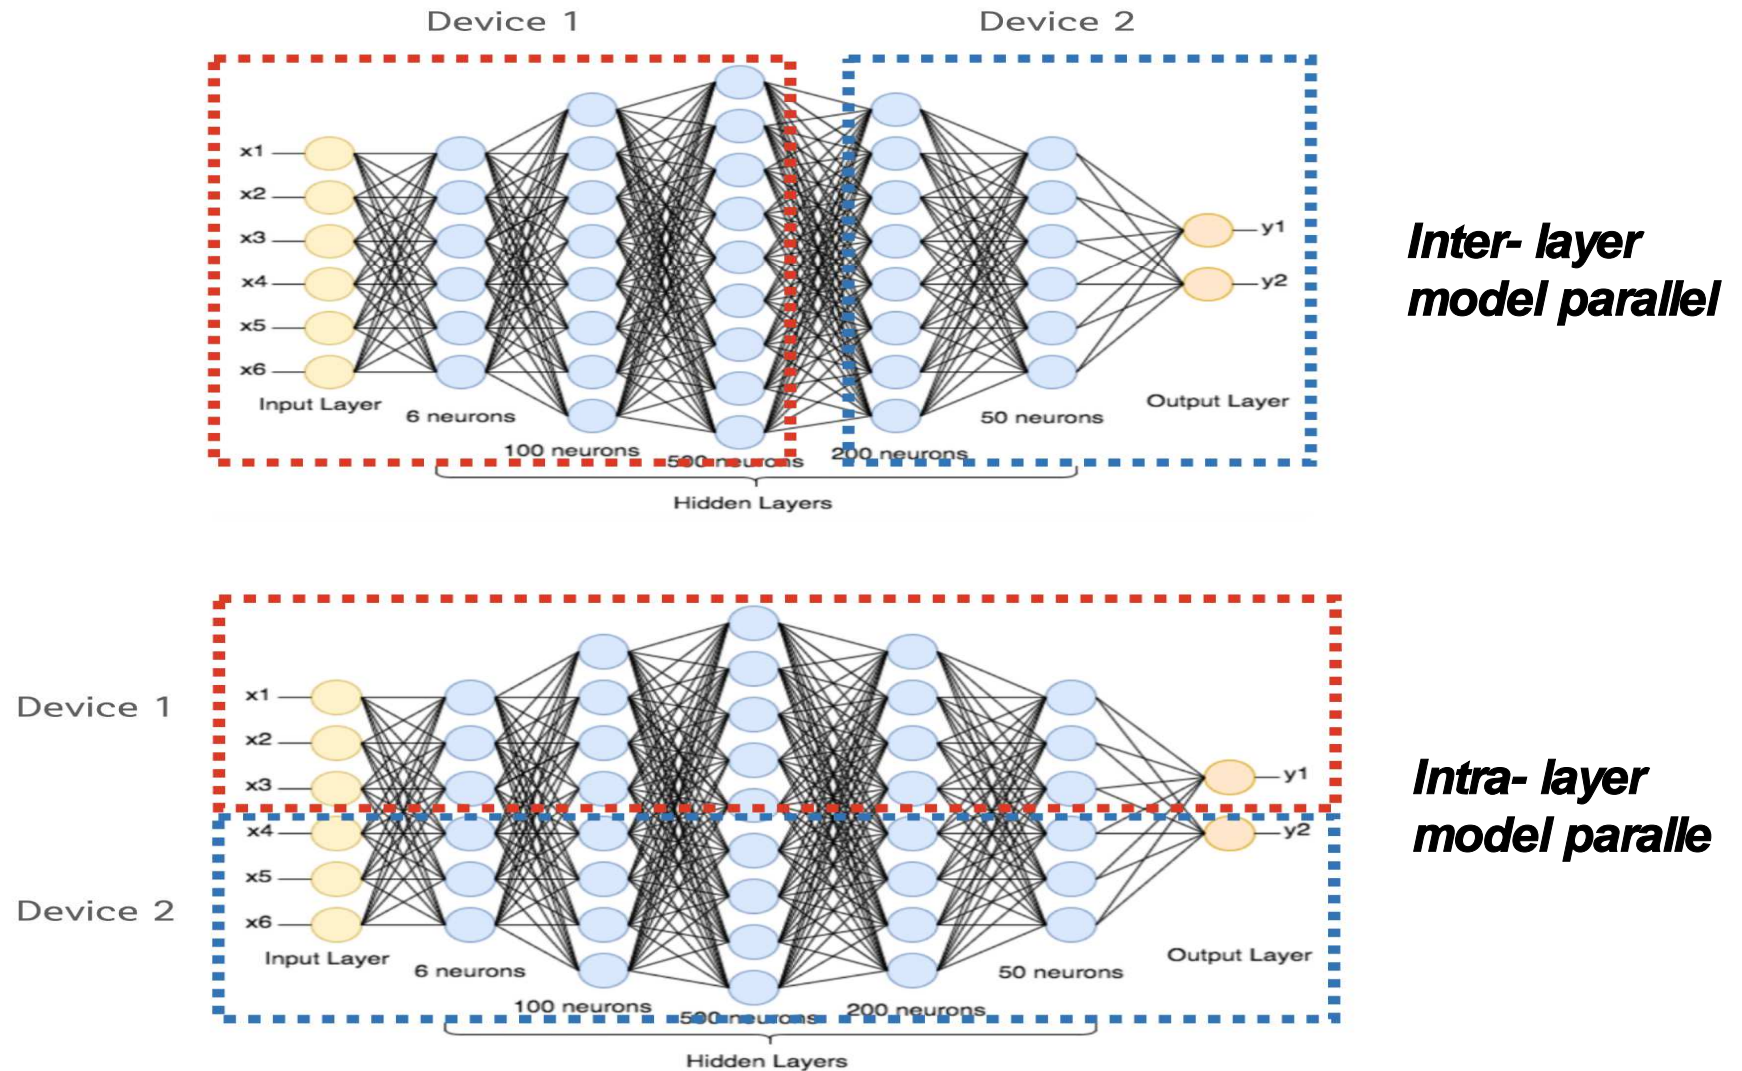
- **Inter-layer Model Parallel**
 - pipeline parallel
- **Intra-layer Model Parallel**
 - Tensor parallel
- **Multi-dimensional Parallel**
 - 3D parallel

Intra- layer parallel




Inter- layer parallel

Inter/Intra-layer Model Parallel



Pipeline Model Parallel



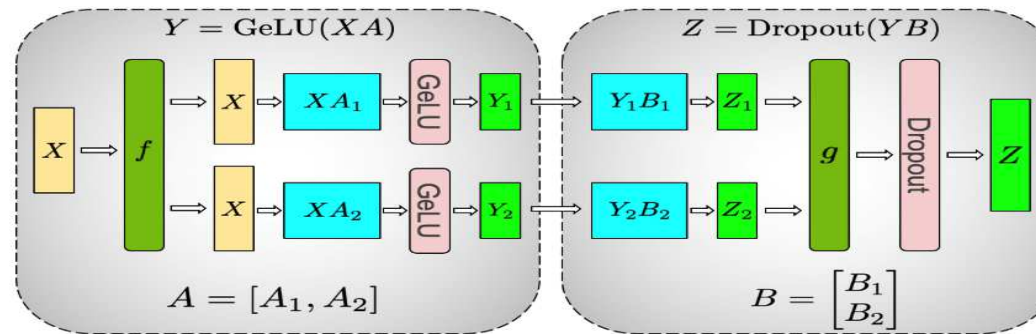
GPipe: Easy Scaling with Micro- Batch Pipeline Parallelism

<https://arxiv.org/pdf/1811.06965.pdf>


PipeDream: Fast and Efficient Pipeline Parallel DNN Training

<https://arxiv.org/pdf/1806.03377.pdf>

Tensor Parallel



(a) MLP



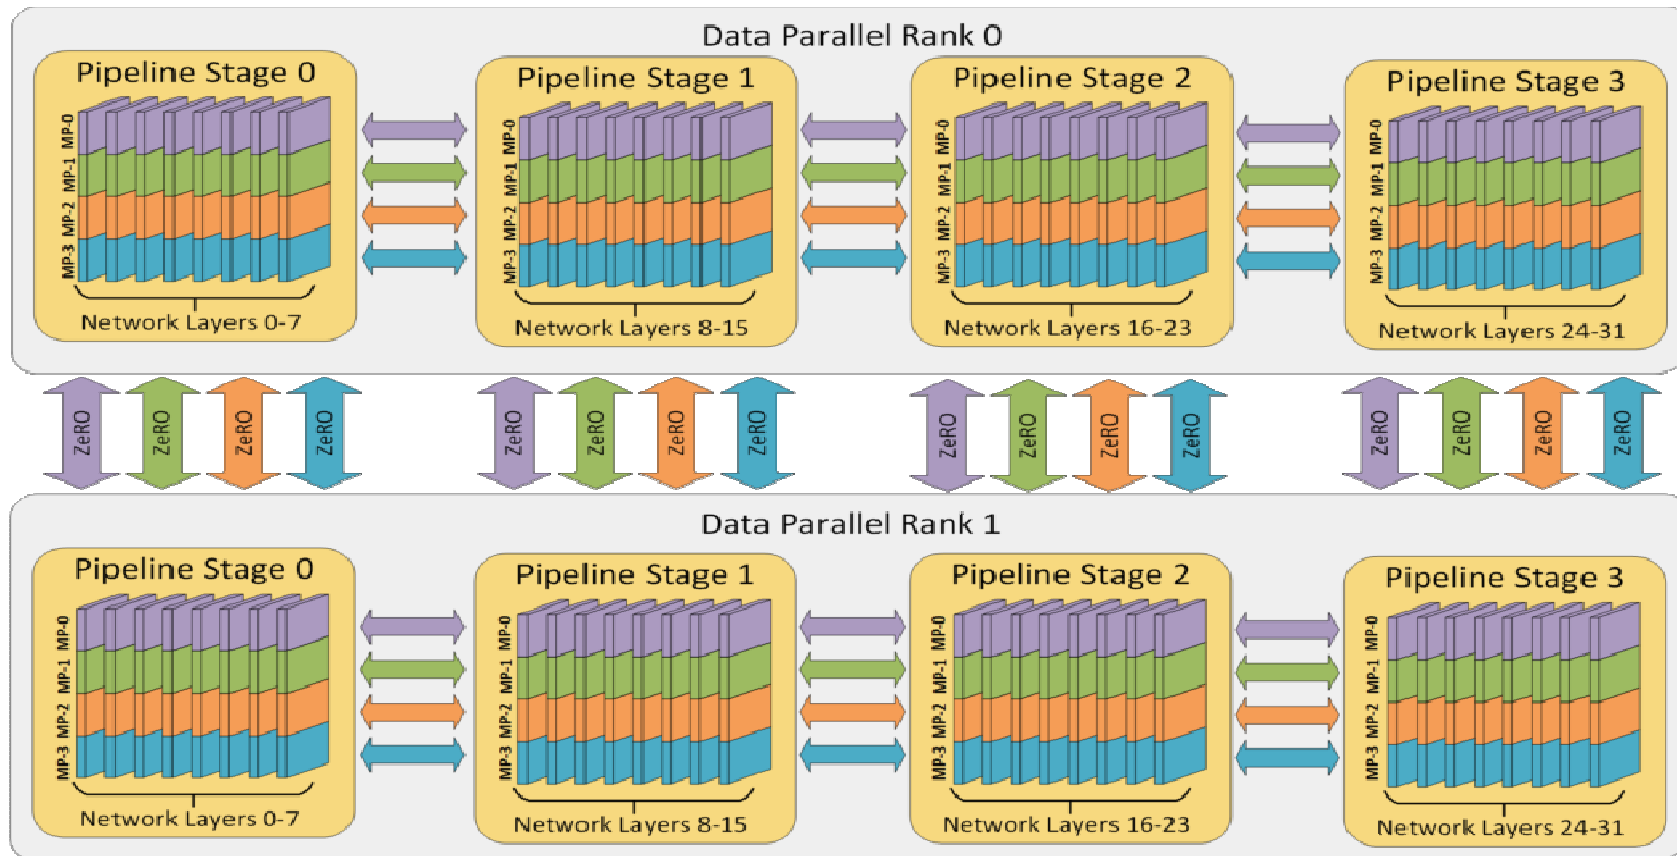
(b) Self attention

Megatron- LM: Training Multi- Billion Parameter Language Models Using Model Parallelism

<https://arxiv.org/pdf/1909.08053.pdf>

3D Parallel

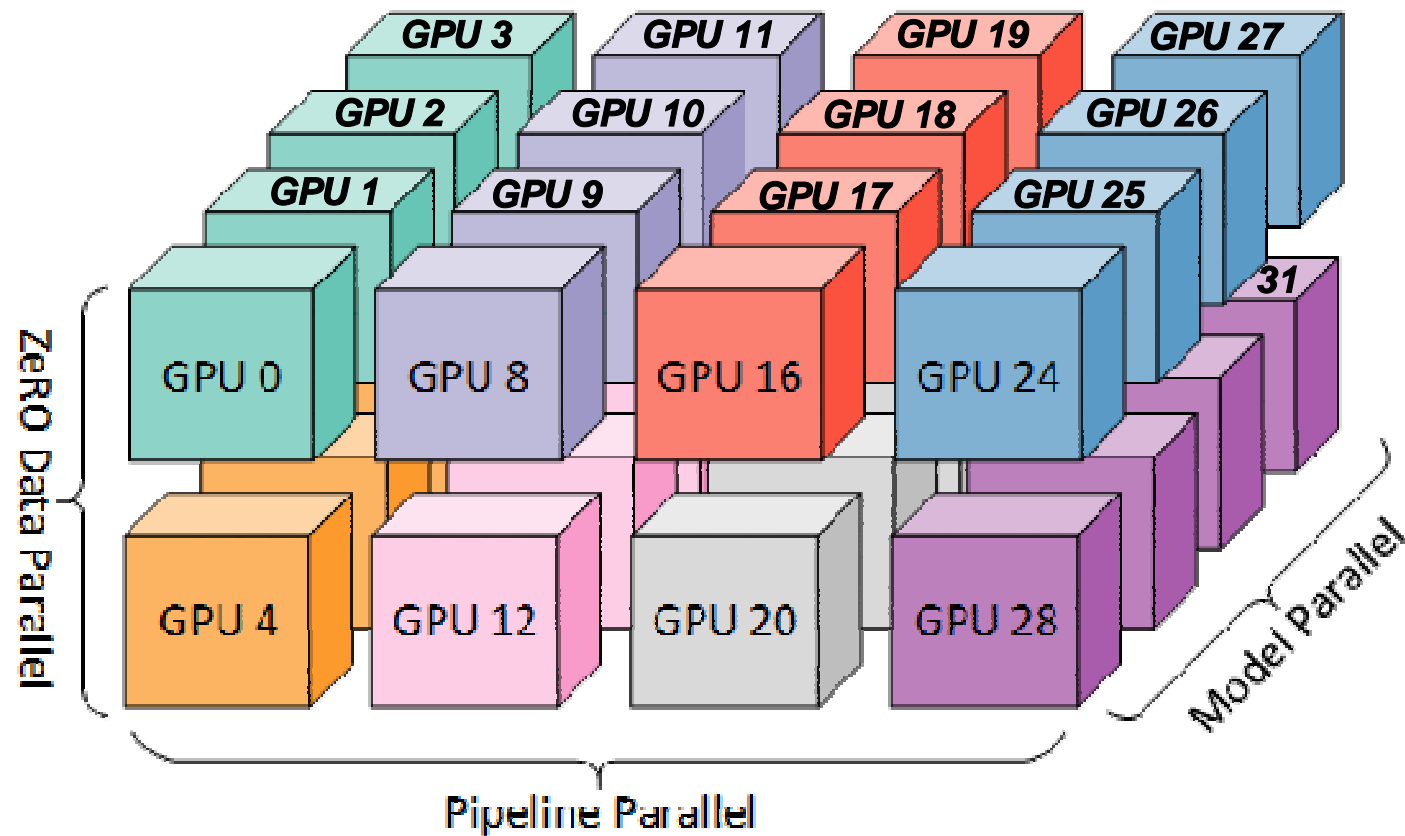
- **Data Parallel + Pipeline Parallel + Tensor Parallel**
 - 32 workers (2DP x 4PP x 4TP)



DeepSpeed: Extreme- scale model training for everyone

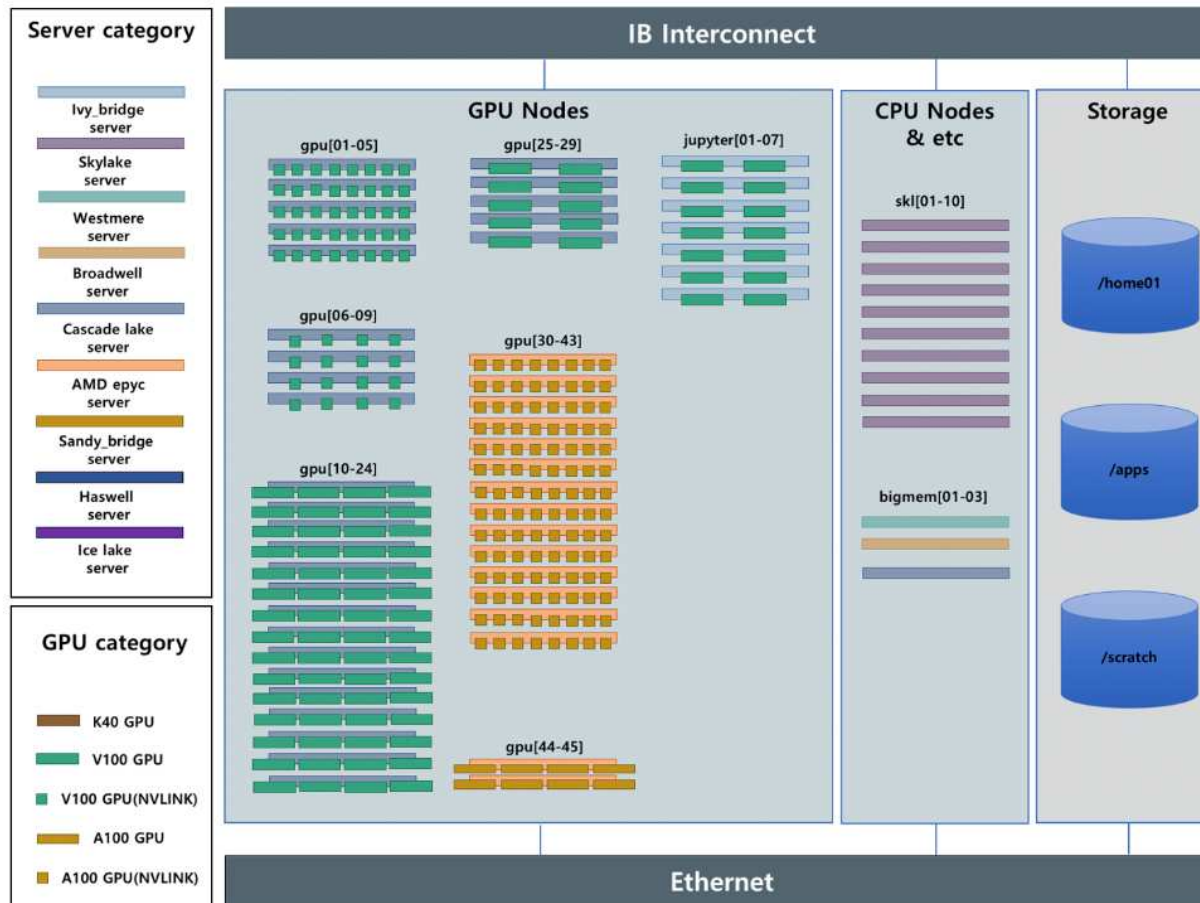
<https://www.microsoft.com/en-us/research/blog/deepspeed-extreme-scale-model-training-for-everyone>

3D Parallel Example (2 x 4 x 4)



KISTI GPU Cluster: Neuron

KISTI GPU Cluster: Neuron



year	#node	#GPU	#Flops
2018	44	53	312.8TF
2019	63	99	698.8TF
2020	78	163	1.24PF
2021	59	200	2.34PF
2022	65	260	3.53PF

Slurm Queues on Neuron

Queue Name (CPU_GPU_GPU#)	#Node	#Total CPU Core	#Job Submission Limit per User	#GPU allocation Limit per User	
cas_v100nv_8	5	160	2	40	V100 (NVlink) 8ea
cas_v100nv_4	4	160	2	16	V100 (NVlink) 4ea
cas_v100_4	15	600	4	40	V100 4ea
cas_v100_2	5	160	2	10	V100 2ea
amd_a100nv_8	14	868	4	64	A100 (Nvlink) 8ea
amd_a100_4	2	128	1	8	A100 4ea
amd_a100_2	2	128	1	4	A100 2ea
skl	10	360	2	-	
bigmem	3	120	1	-	

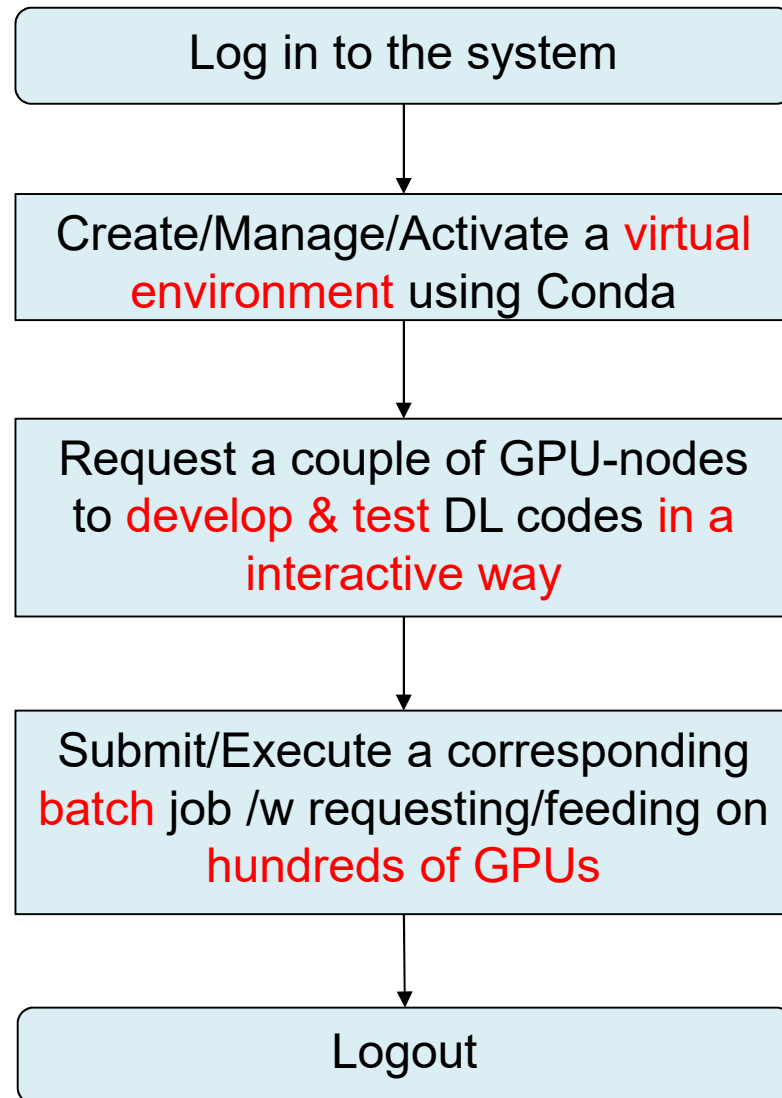
Slurm Queues & available GPUs on Neuron

```
$ sinfo
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST
jupyter   up 2-00:00:00 3 mix jupyter[02-04]
jupyter   up 2-00:00:00 1 idle jupyter01
cas_v100nv_8 up 2-00:00:00 1 mix gpu01
cas_v100nv_8 up 2-00:00:00 4 alloc gpu[02-05]
cas_v100nv_4 up 2-00:00:00 1 mix gpu09
cas_v100nv_4 up 2-00:00:00 2 alloc gpu[07-08]
cas_v100_4 up 2-00:00:00 2 mix gpu[13,17]
cas_v100_4 up 2-00:00:00 10 alloc gpu[10-12,18-24]
cas_v100_4 up 2-00:00:00 2 idle gpu[14-15]
cas_v100_2 up 2-00:00:00 1 mix gpu25
cas_v100_2 up 2-00:00:00 1 alloc gpu26
amd_a100nv_8 up 2-00:00:00 2 mix gpu[36-37]
amd_a100nv_8 up 2-00:00:00 6 alloc gpu[30,32-33,39-41]
amd_a100nv_8 up 2-00:00:00 1 idle gpu42
amd_a100_4 up 2-00:00:00 1 mix gpu44
amd_a100_4 up 2-00:00:00 1 alloc gpu45
skl       up 2-00:00:00 8 idle skl[01-06,08-09]
bigmem    up 2-00:00:00 2 idle bigmem[01-02]
exclusive up infinite 1 mix gpu06
scidebert up infinite 1 mix gpu35
scidebert up infinite 1 alloc gpu34
new_service up infinite 4 down* bigmem03,jupyter[05-06],skl07
maintenance up infinite 6 idle gpu[16,29,31,38,43],jupyter07
```

- **Slurm Quick Start User Guide**
 - <https://slurm.schedmd.com/quickstart.html>
- **KISTI User Guide**
 - <https://www.ksc.re.kr/gsjw/jcs/hd>

Distributed Training Practices on Supercomputer

Distributed Training Practices on Neuron



create/activate a virtual environment

```
$ conda create -n pt_env python=3.7
```

```
$ conda activate pt_env
```

```
$ (pt_env) conda install pytorch
```

```
$ (pt_env) python train.py
```

request 2 nodes for interactive job

```
$ salloc --nodes=2 --time=8:00:00 --  
gres=gpu:4 # available GPU-nodes allocated
```

run& test ML/DL codes interactively

```
$ (pt_env) srun -n 8 python train_ddp.py
```

submit a batch job script requesting 20 nodes with 8 GPUs each node

```
$ (pt_env) sbatch train_ddp_script.sh
```

monitor/check the job status

```
$ (pt_env) queue
```

Conda Virtual Environment

Anaconda vs. PIP

- **Anaconda**

- distribution of the Python and R programming languages for scientific computing
 - data science, machine learning applications, large-scale data processing, predictive analytics, etc
- aims to simplify package management and deployment

- **PIP**


- package installer for Python
- use pip to install packages from the Python Package Index and other indexes.

- **Conda**

- open source **package management** system and **environment management** system
- runs on Windows, macOS, Linux and z/OS
- quickly installs, runs and updates packages and their dependencies.

Miniconda

- **a small, bootstrap version of Anaconda**
 - includes only conda, Python, the packages they depend on
- **a free minimal installer for conda**



Anaconda vs Miniconda

- **Number of packages**
 - Anaconda comes with over 150 data science packages, whereas miniconda comes with only a handful
- **Interface**
 - Anaconda has a graphical user interface (GUI) called the Navigator, while miniconda has a command-line interface

```
### Anaconda Download
```

```
$ wget https://repo.anaconda.com/archive/Anaconda3-2022.10-Linux-x86_64.sh
```

```
### Miniconda Download
```

```
$ wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh
```

Conda command

clean	Remove unused packages and caches.
config	Modify configuration values in .condarc. This is modeled after the git config. Writes to the user .condarc file (/home01/userID/.condarc) by default.
create	Create a new conda environment from a list of specified packages.
help	Displays a list of available conda commands and their help strings.
info	Display information about current conda install.
init	Initialize conda for shell interaction. [Experimental]
install	Installs a list of packages into a specified conda environment.
list	List linked packages in a conda environment.
package	Low-level conda package utility. (EXPERIMENTAL)
remove	Remove a list of packages from a specified conda environment.
uninstall	Alias for conda remove.
run	Run an executable in a conda environment. [Experimental]
search	Search for packages and display associated information. The input is a MatchSpec, a query language for conda packages. See examples below.
update	Updates conda packages to the latest compatible version.
upgrade	Alias for conda update

Conda Virtual Environment

■ Download & Install Conda on /scratch/userID directory

```
## Miniconda
$ wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh
$ chmod 755 Miniconda3-latest-Linux-x86_64.sh
$ ./Miniconda3-latest-Linux-x86_64.sh
```

- ✓ Type your Conda installation directory to **“/scratch/userID/miniconda3”**
 - ✓ Conda default installation directory: \$HOME/miniconda3
- ✓ Type Conda init: **“yes”**, which will add conda init scripts to ~/.bashrc

```
$ source ~/.bashrc # set conda path
$ conda config --set auto_activate_base false
$ which conda
/scratch/$USER/miniconda3/condabin/conda
$ conda --version
conda 4.12.0
$ ls /scratch/$USER/miniconda3
./          conda-meta/ lib/      mkspecs/   qml/       ssl/
../         doc/        libexec/   phrasebooks/ resources/  translations/
bin/        envs/       LICENSE.txt pkgs/       sbin/      var/
....
```

Horovod


What is Horovod?

- **distributed deep learning training framework for TensorFlow, Keras, PyTorch, and Apache MXNet**
 - developed by Uber in 2017
- **aims to make distributed deep learning fast and easy to use**

Why Horovod?

- neutral to DL frameworks to be used
 - ✓ Is it possible to make your DL codes run in parallel, irrespective of whether you use Tensorflow, Keras or Pytorch frameworks?
- easy to use & codify
 - ✓ How much modification does one have to make to a existing DL code to make it distributed?
 - ✓ How easy is it to run it?
- fast to run
 - How much faster would it run in distributed mode?
 - how easy is it to scale up?

Scaling with Horovod



Horovod Usage

- **5 steps/lines to be added in your code**
 - Initialize Horovod
 - Pin GPU to each worker
 - Wrap the optimizer
 - Synchroize state across workers
 - Checkpoint on the first worker

Initialize Horovod

- **Tensorflow**
 - `import horovod.tensorflow as hvd`
 - `hvd.init()`
- **Kera**
 - `import horovod.keras as hvd`
 - `hvd.init()`
- **Pytorch**
 - `import horovod.torch as hvd`
 - `hvd.init()`

Pin a GPU for each worker

- **Tensorflow**

- `tf.config.experimental.set_visible_devices(gpus[hvd.local_rank()], 'GPU')`

- **Keras**

- `tf.config.experimental.set_visible_devices(gpus[hvd.local_rank()], 'GPU')`

- **Pytorch**

- `torch.cuda.set_device(hvd.local_rank())`

Adjust learning rate and wrap the optimizer

- **Tensorflow**

- `opt = tf.optimizers.Adam(0.01 * hvd.size())`
- `opt = hvd.DistributedOptimizer(opt,...)`

- **Keras**

- `opt = keras.optimizers.Adadelta(0.01 * hvd.size())`
- `opt = hvd.DistributedOptimizer(opt,...)`

- **Pytorch**

- `opt = optim.SGD(model.parameters(), 0.01 * hvd.size())`
- `opt= hvd.DistributedOptimizer(opt, ...)`

Synchronize states across workers

- **Tensorflow/Kera**
 - `callbacks =`
`[hvd.callbacks.BroadcastGlobalVariablesCallback(0)]`
- **Pytorch**
 - `hvd.broadcast_parameters(model.state_dict(),`
`root_rank=0)`
 - `hvd.broadcast_optimizer_state(optimizer, root_rank=0)`
- **Ensure all workers start with the same weights**

Checkpoint on the first worker (rank 0)

- **Tensorflow/Keras**

- if `hvd.rank() == 0`:
 `callbacks.append(keras.callbacks.ModelCheckpoint(args.c
 heckpoint_format))`

- **Pytorch**

- if `hvd.rank() == 0`:
 `state = {'model': model.state_dict(),
 'optimizer': optimizer.state_dict(), }`
 `torch.save(state, filepath)`

Pytorch Example

```
import torch
import horovod.torch as hvd

# Initialize Horovod
hvd.init()

# Horovod: pin GPU to local rank.
torch.cuda.set_device(hvd.local_rank())

# Build model.
model = Net()
model.cuda()
optimizer = optim.SGD(model.parameters())

# Horovod: wrap optimizer with DistributedOptimizer.
optimizer = hvd.DistributedOptimizer(
    optimizer,
    named_parameters=model.named_parameters())

# Horovod: broadcast parameters.
hvd.broadcast_parameters(
    model.state_dict(),
    root_rank=0)

for epoch in range(100):
    for batch_idx, (data, target) in enumerate(...):
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)
        loss.backward()
        optimizer.step()
```


Horovod execution command

- **MPI takes care of launching processes on all nodes**
- **Run on a 4-GPU machine**
 - `$ salloc ...`
 - `$ mpirun -np 4 -H localhost:4 python train_hvd.py`
 - `$ horovodrun -np 4 -H localhost:4 python train_hvd.py`
 - `$ srun -n python train_hvd.py`
- **Run on 4 machines with 4-GPUs each**
 - `$ salloc ...`
 - `$ mpirun/horovodrun -np 4 -H
node1:4,node2:4,node3:4,node4:4 python train_hvd.py`
 - `$ srun -n 16 python train_hvd.py`

Horovod Installation on Neuron

```
$ module load gcc/10.2.0 cuda/11.4 cudamp/openmpi-4.1.1 cmake/3.16.9
$ conda create -n horovod
$ conda activate horovod
$ conda install pytorch==1.12.0 torchvision==0.13.0 torchaudio==0.12.0 cudatoolkit=11.3 -c pytorch
$ pip install tensorflow-gpu==2.10.0
$ HOROVOD_GPU_OPERATIONS=NCCL HOROVOD_WITH_TENSORFLOW=1
HOROVOD_WITH_PYTORCH=1 \
HOROVOD_WITH_MPI=1 HOROVOD_WITH_GLOO=1 pip install --no-cache-dir horovod
$ horovodrun -cb
Horovod v0.26.1:
```

Available Frameworks:

```
[X] TensorFlow
[X] PyTorch
[ ] MXNet
```

Available Controllers:

```
[X] MPI
[X] Gloo
```

Available Tensor Operations:

```
[X] NCCL
[ ] DDL
[ ] CCL
[X] MPI
[X] Gloo
```

Horovod execution on Neuron

1) request an allocation of available on neuron

```
$ salloc --partition=amd_a100nv_8 -J debug --nodes=2 --time=2:00:00 --gres=gpu:4 --comment=horovod
```

2) load modules

```
$ module load gcc/10.2.0 cuda/11.4 cudamp/openmpi-4.1.1 cmake/3.16.9
```

3) activate the horovod virtual environment

```
$ conda activate horovod
```

4) run horovod applications

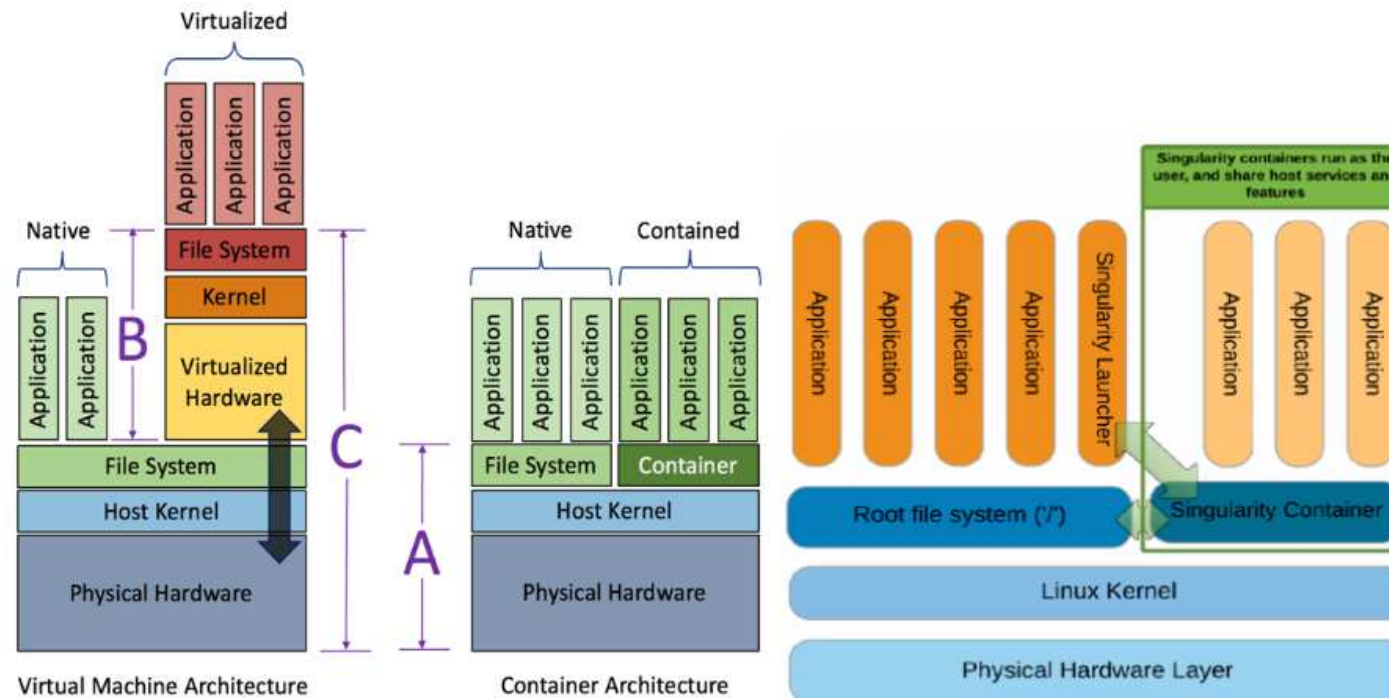
```
(horovod) $ mpirun/horovod -np 8 -H gpu#:4,gpu#:4 python train_hvd.py
```

```
(horovod) $ srun -n python train_hvd.py
```

Singularity Container

Singularity

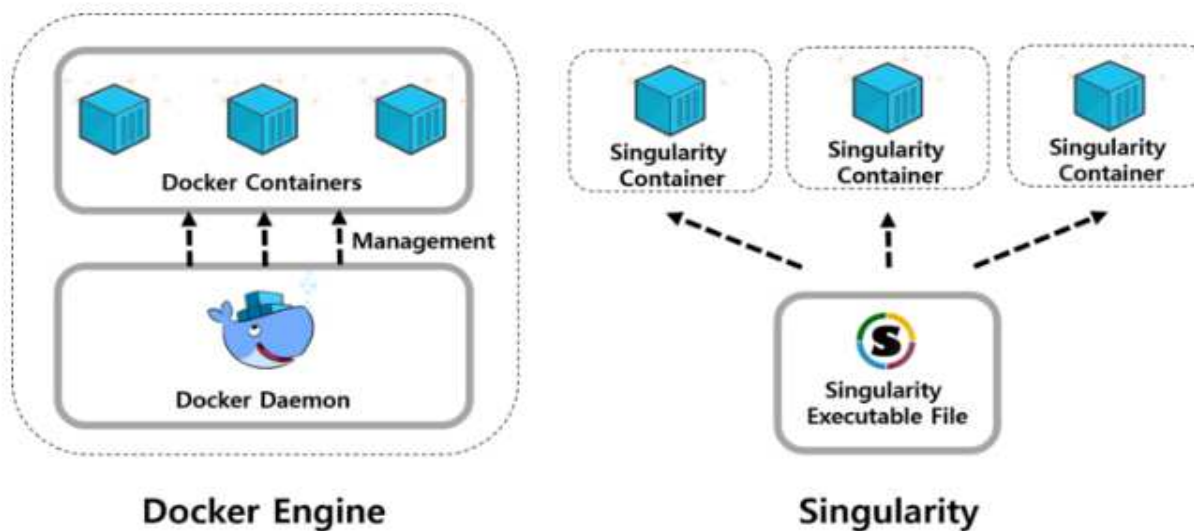
- a container platform for HPC



Container-based applications have *direct access* to the host kernel and hardware and, thus, are able to achieve similar performance to native applications. In contrast, VM-based applications only have *indirect access* via the guest OS and hypervisor, which creates a significant performance overhead.

Why Singularity?

- **A container platform for HPC**
 - Each container is a single image file
 - No root owned daemon processes
 - Support shared/multi-tenant resource environment
 - Support HPC hardware
 - Infiniband, GPUs
 - Support HPC applications
 - MPI



Running Horovod using Singularity on Neuron

- No bother to deal with conda & horovod
- Just allocate nodes using salloc and run singularity container. **That's it!!!**

```
$ salloc --partition=amd_a100nv_8 -J debug --nodes=2 --time=2:00:00 --gres=gpu:4 --comment=pytorch
```

```
$ srun -n 8 singularity exec --nv  
/apps/applications/singularity_images/ngc/pytorch_22.03-hd-py3.sif python  
pytorch_imageNet_resnet50.py
```

```
/usr/bin/rm: cannot remove '/usr/local/cuda/compat/lib': Read-only file system
```

```
/usr/bin/rm: cannot remove '/usr/local/cuda/compat/lib': Read-only file system
```

```
.....
```

```
Train Epoch   #10:  0%|          | 1/5005 [00:13<19:03:52, 13.72s/it, loss=2.29, accurTrain Epoch  
#10:  0%|          | 1/5005 [00:13<19:03:52, 13.72s/it, loss=2.22, accurTrain Epoch   #10:  0%|          |  
2/5005 [00:13<19:03:38, 13.72s/it, loss=2.16, accurTrain Epoch   #10:  0%|          | 3/5005  
[00:13<5:00:52, 3.61s/it, loss=2.16, accurTrain Epoch   #10:  0%|          | 3/5005 [00:13<5:00:52,  
3.61s/it, loss=2.17, accurTrain Epoch   #10:  0%|          | 4/5005 [00:13<5:00:48, 3.61s/it, loss=2.23,  
accura
```

```
.....
```

Singularity Usage on Neuron

- **Web site:** <https://www.ksc.re.kr/gsjw/jcs/hd>
- **job script directory**
 - /apps/applications/singularity_images/examples
- **Singularity Container Images directory**
 - /apps/applications/singularity_images/ngc
- **Pytorch examples directory**
 - Single node
 - /apps/applications/singularity_images/examples/pytorch/resnet50v1.5
 - Multiple nodes
 - /apps/applications/singularity_images/examples/horovod/examples/pytorch
- **Imagenet datasets directory**
 - Training data
 - /apps/applications/singularity_images/imagenet/train
 - Validation data
 - /apps/applications/singularity_images/imagenet/val

Thank you

Contact: Soonwook Hwang <hwang@kisti.re.kr>

github.com/hwang2006/KISTI-DL-tutorial-using-horovod